

MCTensor: A High-Precision Deep Learning Library with Multi-Component Floating-Point

Tao Yu*, Wentao Guo*, Jianan Canal Li*, Tiancheng Yuan*, Christopher De Sa

*Equal Contribution

Department of Computer Science, Biological Engineering, and System Engineering

Cornell University



Cornell Bowers C-IS

College of Computing and Information Science

High Precision Computations

Applications: dynamical systems (Taylor methods) [1], computational geometry (delaunay triangulation) [2], hyperbolic deep learning [3,4] ...

How to achieve high precision arithmetic?

- multiple-digit: a sequence of digits coupled with a single exponent, e.g., GNU Multiple Precision (GMP) Arithmetic Library [5], Julia's BigFloat type [6] ...
- multiple-component (MCF) [7]: an unevaluated sum of multiple ordinary floating-point numbers (e.g., float16, float32)

What's the performance difference?

- multiple-digit can represent compactly a much larger range of numbers;
- multiple-component can use existing floating-point accelerators and hence faster.



Outline

- MCTensor Library
- Computing & Learning with MCTensor
- Error Analysis of MCTensor
- Experiment Evaluation
- Conclusions and Future Work



MCTensor Library

- MCF as underlying representation using an “expansion”:

$$x = (x_0, x_1, \dots, x_{nc-1}) = x_0 + x_1 + \dots + x_{nc-1}$$

where nc is the number of components and each component x_i is an ordinary floating-point (e.g., PyTorch FloatTensor).

- MCTensor Object

$$x\{fc, \text{tensor}, nc\}$$

- Gradient of a MCTensor

$$\frac{\partial f}{\partial x} = \frac{\partial f}{\partial (x_0 + x_1 + \dots + x_{nc-1})} = \frac{\partial f}{\partial x_i}$$



Computing with MCTensor

Basic operators based on inputs:

- **MCTensor** and **Tensor**: Grow-ExpN (add), ScalingN (multiply), DivN (divide)
- **MCTensor** and **MCTensor**: Add-MCN (add), Mul-MCN (multiply), Div-MCN (divide)
- **MCTensor**: Exp-MCN (exp), Square-MCN (square)

where **Two-Sum(Tensor1, Tensor2)** \rightarrow (result, error), **Simple-Renorm(h, nc)** \rightarrow moves zeros backward and outputs a MCTensor with nc component.


Matrix Operators

- **AddMM-MCN** (torch.addmm)
- **MV-MCN** (torch.mv)
- **Matmul-MCN** (torch.matmul)
- ...

We extend and broadcast various computations supported for PyTorch Tensor to MCTensor! Compatible with PyTorch, and easy to plug in.

Algorithm 1 Grow-ExpN

Input: nc -MCTensor x , PyTorch Tensor v
initialize $Q \leftarrow v$
for $i = 1$ to nc **do**
 $k \leftarrow nc + 1 - i$
 $(Q, h_k) \leftarrow \text{Two-Sum}(x_{k-1}, Q)$
end for
 $h \leftarrow (Q, h_1, \dots, h_{nc})$
Return: **Simple-Renorm**(h, nc)



Learning with MCTensor

MCMModule: e.g., MCLinear, MCEmbedding, MCSequential ...

MCActivation: e.g., MCSoftmax, MCReLU, MC-GELU ...

MCOptimizer: e.g., MC-SGD, MCAdam ...

```
class MCLinear(MCMModule):
    def __init__(self, in_features, out_features, nc, bias=True):
        super(MCLinear, self).__init__()
        self.in_features = in_features
        self.out_features = out_features
        self.nc = nc
        self.weight = MCTensor(out_features, in_features, nc=nc, requires_grad=True)
        if bias:
            self.bias = MCTensor(out_features, nc=nc, requires_grad=True)
        else:
            self.bias = None

    def forward(self, input):
        return F.linear(input, self.weight, self.bias)
```


```
model = MCModel(nc=nc)
optimizer = MCSGD(mc_model.parameters())

for x, y in train_dataset:
    optimizer.zero_grad()
    y_hat = model(x)
    loss = loss_fn(y_hat, y)
    loss.backward()
    optimizer.step()
```

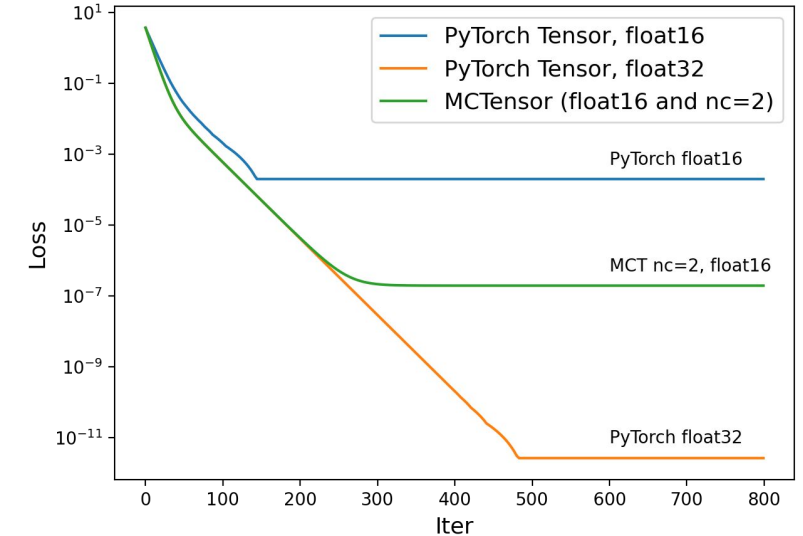

MCTensor can be used in the same way as PyTorch Tensor with a few lines of replacement!



Error Analysis



Relative error of MCTensor computation
between values x, y at different magnitudes




Training loss of linear regression
task on a synthetic dataset

Achieve arbitrary precision computations/training by simply increasing the number of components!



Experiments

- (1) high precision learning with low precision numbers, using Logistic Regression and MLP model;
- (2) hyperbolic embedding using the Poincaré Halfspace model [8].



Logistic Regression on Breast Cancer Dataset

MCTensor demonstrated better precision and performance in both settings!

Model	Training Loss	Testing accuracy
MLP Float16	0.144	90.35
MLP Float32	0.124	91.23
MLP Float64	0.124	91.23

MC-MLP (nc=1)	0.144	90.35
MC-MLP (nc=2)	0.124	91.23
MC-MLP (nc=3)	0.124	91.23

MLP on Breast Cancer Dataset

Model	MAP (mean \pm sd)	MR (mean \pm sd)
Halfspace (f32)	91.91% \pm 0.64%	1.399 \pm 0.04
Halfspace (f64)	92.79% \pm 0.41%	1.340 \pm 0.07

MC-Halfspace (f64 nc=1)	93.02% \pm 0.40%	1.296 \pm 0.02
MC-Halfspace (f64 nc=2)	92.77% \pm 0.28%	1.304 \pm 0.02
MC-Halfspace (f64 nc=3)	93.31% \pm 0.75%	1.282 \pm 0.03

Performance of Hyperbolic Models

MAP \uparrow : mean average precision, MR \downarrow : mean rank



Conclusions and Future Work

- MCTensor achieves high-precision arithmetic while leveraging the benefits of heavily-optimized ordinary floating-point arithmetic
- MCTensor is easy to use: same as PyTorch code with a few lines of replacement
- MCTensor achieves better precision using low precision numbers
- MCTensor helps relieve the imprecision problem in hyperbolic learning
- A promising future work: design and optimize MCTensor for better efficiency on general tasks



Thank you!
Questions & Comments



References

- [1] Bailey et al. High-precision arithmetic in mathematical physics. *Mathematics*, 3(2):337–367, 2015.
- [2] Schirra, S. Robustness and precision issues in geometric computation. 1998.
- [3] Yu et al. Numerically accurate hyperbolic embeddings using tiling-based models. *Advances in Neural Information Processing Systems*, 32, 2019.
- [4] Yu et al. Representing hyperbolic space accurately using multi-component floats. *Advances in Neural Information Processing Systems*, 34, 2021.
- [5] Granlund et al. GNU MP: The GNU Multiple Precision Arithmetic Library, 5.0.5 edition, 2012. <http://gmplib.org/>.
- [6] Bezanson et al. Julia: A fresh approach to numerical computing. *SIAM review*, 59(1):65–98, 2017
- [7] Priest, D. M. Algorithms for arbitrary precision floating point arithmetic. University of California, Berkeley, 1991
- [8] Nickel et al. Poincaré embeddings for learning hierarchical representations. *Advances in Neural Information Processing Systems*, 30, 2017



Operators	Inputs sizes	FloatTensor	1-MCTensor	2-MCTensor	3-MCTensor
Dot-MCN	5000, 5000	$1.61\mu s \pm 3.29ns$	$442\mu s \pm 5.61\mu s$	$656\mu s \pm 1.16\mu s$	$858\mu s \pm 12.2\mu s$
MV-MCN	$(5000 \times 500), 500$	$157\mu s \pm 4.32\mu s$	$320ms \pm 5.78ms$	$460ms \pm 10.7ms$	$580ms \pm 12.1ms$
Matmul-MCN	$(500 \times 200), (200 \times 50)$	$97.3\mu s \pm 1.1\mu s$	$495ms \pm 10.8ms$	$735ms \pm 21.7ms$	$934ms \pm 28ms$

Operators	Inputs sizes	FloatTensor	1-MCTensor	2-MCTensor	3-MCTensor
Add-MCN	(1000×1000)	$497\mu s \pm 6.77\mu s$	$26.7ms \pm 486\mu s$	$44.7ms \pm 379\mu s$	$64.4ms \pm 385\mu s$
ScalingN	(1000×1000)	$490\mu s \pm 9.69\mu s$	$33.7ms \pm 402\mu s$	$57.5ms \pm 842\mu s$	$84.7ms \pm 1.78ms$
Mult-MCN	(1000×1000)	$514\mu s \pm 15.2\mu s$	$218ms \pm 4.03ms$	$667ms \pm 11.6ms$	$1.4s \pm 21.6ms$
Div-MCN	(1000×1000)	$510\mu s \pm 10.6\mu s$	$80.3ms \pm 770\mu s$	$243ms \pm 3.24ms$	$498ms \pm 8.26ms$

Table 4. MCTensor Basic Operators Running Time (mean \pm sd)

$$d_u(\mathbf{x}, \mathbf{y}) = \text{arcosh}\left(1 + \frac{\|\mathbf{x} - \mathbf{y}\|^2}{2x_n y_n}\right)$$

$$\mathcal{L}(\Theta) = \sum_{(\mathbf{x}, \mathbf{y}) \in D} \log \frac{e^{-d_u(\mathbf{x}, \mathbf{y})}}{\sum_{\mathbf{y}' \in \mathcal{N}(\mathbf{x})} e^{-d_u(\mathbf{x}, \mathbf{y}')}}}$$